

LASYA PONNAPATI

SOFTWARE ENGINEER / AI ENGINEER

lasyaponnapati@gmail.com | +91 8309658235 | Bangalore, India | [LinkedIn](#) | [GitHub](#)

EXPERIENCE

Software Engineer — Asama.ai

July 2025 – Present

Bengaluru, India

- Developed **multi-agent workflows** using **CrewAI**, implementing **agent roles**, **custom tools**, and **orchestration logic** to support infrastructure monitoring and troubleshooting across **500+ servers**.
- Engineered and iteratively refined **10+ LLM prompts** and **intent-based routing** logic across **5+ infrastructure data sources**, improving query routing and agent response relevance by **30%**.
- Integrated **LangChain** for LLM orchestration and **Langfuse** for observability, enabling end-to-end prompt tracing and performance monitoring across **1,000+ agent executions**, reducing LLM workflow debugging time by **40%**.
- Implemented real-time **LLM response streaming** using **Server-Sent Events (SSE)** through **backend event pipelines** and **frontend rendering**, reducing perceived response latency by **50%** through incremental response delivery.
- Developed **10+ storage metric collectors** for an infrastructure exporter, enabling **Prometheus/VictoriaMetrics** ingestion and expanding storage health telemetry coverage across **500+ infrastructure nodes**.
- Designed and implemented a **Neo4j** infrastructure knowledge graph using **Cypher** queries, modeling **500+ infrastructure nodes** and **20+ relationship types** to power product dashboards and provide **graph-based retrieval** for AI-powered natural language queries.
- Improved **DBSCAN/HDBSCAN**-based infrastructure anomaly detection by refining **hardware-aware server clustering**, reducing false-positive anomalies by **50%** and improving detection reliability across **500+ monitored servers**.
- Enhanced **rule-based** infrastructure **fault detection** by refining **25+ health metric evaluation rules** across **10+ server components**, reducing false-positive fault detections by **30%**.
- Built a threshold-based infrastructure alerting system using **vmaalert** and **Alertmanager**, supporting **100+ alert rules** across **500+ monitored servers**, and developed a **notification service** routing tenant-configured alerts across 4 channels: **Slack**, **Email**, **Mattermost**, and **Google Chat**.

Software Engineer Intern — SafeSend Technologies

Feb 2024 – Sept 2024

Bengaluru, India

- Developed **RESTful APIs** in **C#/ASP.NET** and functional **React components**, improving front-end/back-end integration and codebase maintainability.
- Implemented **backend unit tests** and designed **SQL stored procedures** to **optimize data retrieval** and query performance.

SKILLS

AI and LLM Engineering: Generative AI, LLMs, Prompt Engineering, Context Engineering, LLM Orchestration, AI Agents, Multi-Agent Systems, Semantic Search, Information Retrieval, RAG, CrewAI, LangChain, LlamaIndex

Languages: Python, Go, JavaScript

Backend & Infrastructure: REST APIs, Microservices, Temporal, Workflow Orchestration, Docker, AWS, Linux

Databases & Graphs: SQL, PostgreSQL, Vector Databases, Neo4j, Knowledge Graphs, etcd, Key-Value Stores

Observability & Monitoring: Langfuse, AI Observability, Prometheus, Telemetry Exporters, Infrastructure Monitoring

Developer Tools: Git, GitHub, Cursor, Claude Code, Codex

CERTIFICATIONS

- Google AI Essentials Specialization**
- Google Prompting Essentials Specialization**

August 2026

August 2026

EDUCATION

Bachelor of Technology, Information Technology —
REVA University

June 2024

Bengaluru, India